

```
In [50]: import requests
        from bs4 import BeautifulSoup

In [51]: user_agent = 'Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/116.0.0.0 Safari/537.36'
        Link_url = 'https://www.imdb.com/chart/top'
        response = requests.get(Link_url, headers = {'User-Agent':user_agent, 'Accept_language':'en-US, en;q=0.5'})
        print(response)

<Response [200]>

In [52]: soup = BeautifulSoup(response.content, 'html.parser')
        movie_data = soup.findAll('li',attrs = {'class':'ipc-metadata-list-summary-item sc-59b6048d-0 jemTre cli-parent'})
        movie_rank = []
        movie_name = []
        released_year = []
        duration = []
        rated_type = []
        ratings = []
        votes = []

In [53]: # To convert votes count from Millions and thousands to numbers
        def convert_to_int (vote):
            if 'm' in vote:
                return int(float(vote.replace('m','')) * 1000000)
            elif 'k' in vote:
                return int(float(vote.replace('k','')) * 1000)
            else:
                return int(vote)

        # Convert duration into minutes
        def convert_to_min(r_time):
            total_minutes = 0

            parts = r_time.split()
            for part in parts:
                if 'h' in part:
                    total_minutes += int(part.replace('h','')) * 60
                elif 'g' in part:
                    total_minutes += int(part.replace('g',''))
            return total_minutes

        for movie in movie_data:
            # Name and rank
            rank_name = movie.find('div',class_="ipc-metadata-list-summary-item__c").div.a.text
            rank = int(rank_name.split('. ')[0])
            name = rank_name.split('. ')[1]
            movie_rank.append(rank)
            movie_name.append(name)

            # year
            year = int(movie.find('div',class_="ipc-metadata-list-summary-item__c").find('span').find_next('span').text)
            released_year.append(year)

            # Duration of movie
            runtime = movie.find('div',class_="ipc-metadata-list-summary-item__c").findAll('span')[2].text
            runtime = convert_to_minutes(runtime)
            duration.append(runtime)

            # Rated Type
            r_type = movie.find('div',class_="ipc-metadata-list-summary-item__tc").find('span').find_next('span').find_next('span').find_next('span')
            rated_type.append(r_type)

            # Ratings
            rating = movie.find('div',class_="ipc-metadata-list-summary-item__c").find('span',class_='ipc-rating-star ipc-rating-star--base ipc-rating-star')
            rating = float(rating)
            ratings.append(rating)

            # votes
            vote = movie.find('div',class_="ipc-metadata-list-summary-item__c").find('span',class_='ipc-rating-star ipc-rating-star--base ipc-rating-star')
            vote = vote.replace('(', '').replace(')', '')
            vote = convert_votes_to_int(vote)
            votes.append(vote)

In [54]: import pandas as pd
        import numpy as np
        df = pd.DataFrame({"Rank": movie_rank, 'Movie_Name' : movie_name,'Duration(Minutes)' : duration, 'Released_Year' : released_year,'Certification' : rated_type})
        df.head(10)

Out[54]:
```

	Rank	Movie_Name	Duration(Minutes)	Released_Year	Certification	Ratings	Votes
0	1	The Shawshank Redemption	142	1994	A	9.3	2800000
1	2	The Godfather	175	1972	A	9.2	2000000
2	3	The Dark Knight	152	2008	UA	9.0	2800000
3	4	The Godfather: Part II	202	1974	A	9.0	1300000
4	5	12 Angry Men	96	1957	U	9.0	834000
5	6	Schindler's List	195	1993	A	9.0	1400000
6	7	The Lord of the Rings: The Return of the King	201	2003	U	9.0	1900000
7	8	Pulp Fiction	154	1994	A	8.9	2200000
8	9	The Lord of the Rings: The Fellowship of the Ring	178	2001	U	8.8	1900000
9	10	Il Buono, Il Brutto, Il Cattivo	161	1966	A	8.8	791000

```
In [55]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 250 entries, 0 to 249
Data columns (total 7 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   Rank                  250 non-null   int64
1   Movie_Name            250 non-null   object
2   Duration(Minutes)     250 non-null   int64
3   Released_Year         250 non-null   int64
4   Certification         250 non-null   object
5   Ratings               250 non-null   float64
6   Votes                 250 non-null   int64
dtypes: float64(1), int64(4), object(2)
memory usage: 13.8+ KB
```

```
In [56]: df.shape
Out[56]: (250, 7)
```

```
In [57]: df.size
Out[57]: 1750
```

```
In [58]: df.ndim
Out[58]: 2
```

```
In [59]: df.describe()
Out[59]:
```

	Rank	Duration(Minutes)	Released_Year	Ratings	Votes
count	250.000000	250.000000	250.000000	250.000000	2.500000e+02
mean	125.500000	129.012000	1986.816000	8.307600	6.719320e+05
std	72.312977	29.756236	25.387086	0.232462	5.418230e+05
min	1.000000	45.000000	1921.000000	8.000000	3.600000e+04
25%	63.250000	107.250000	1966.250000	8.100000	2.295000e+05
50%	125.500000	126.500000	1994.500000	8.200000	5.370000e+05
75%	187.750000	145.750000	2007.000000	8.400000	9.910000e+05
max	250.000000	238.000000	2023.000000	9.300000	2.800000e+06

```
In [60]: top_10_movies = df.sort_values(by='Votes', ascending=False).head(10)

# Add a new column 'Rank' for displaying the rank of each movie
top_10_movies['Rank'] = np.arange(1, len(top_10_movies) + 1)

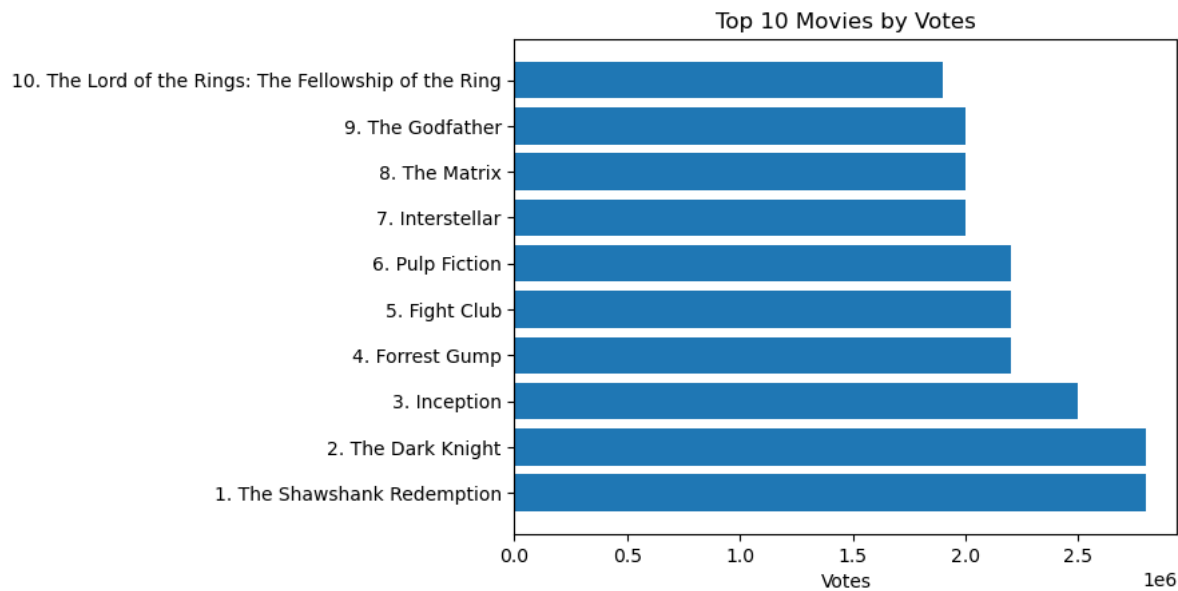
# Display the top 10 movies with their rank, name, and votes
print(top_10_movies[['Rank', 'Movie_Name', 'Votes']])

# Extract Movie Names and Votes
Movie_Name = top_10_movies['Movie_Name']
Votes = top_10_movies['Votes']
Rank = top_10_movies['Rank']

y_pos = np.arange(len(Movie_Name))
plt.barh(y_pos[:-1], Votes[:-1], align='center') # Reverse the order
plt.yticks(y_pos[:-1], [f'{r}. {m}' for r, m in zip(Rank[:-1], Movie_Name[:-1])]) # Reverse the order
plt.xlabel('Votes')
plt.title('Top 10 Movies by Votes')

plt.show()
```

	Rank	Movie_Name	Votes
0	1	The Shawshank Redemption	2800000
2	2	The Dark Knight	2800000
13	3	Inception	2500000
10	4	Forrest Gump	2200000
11	5	Fight Club	2200000
7	6	Pulp Fiction	2200000
22	7	Interstellar	2000000
15	8	The Matrix	2000000
1	9	The Godfather	2000000
8	10	The Lord of the Rings: The Fellowship of the Ring	1900000



```
In [63]: top_10_duration = df.sort_values(by='Duration(Minutes)', ascending=False).head(10)

# Add a new column 'Rank' for displaying the rank of each movie
top_10_duration['Rank'] = np.arange(1, len(top_10_duration) + 1)

# Display the top 10 movies with their rank, name, and duration
print(top_10_duration[['Rank', 'Movie_Name', 'Duration(Minutes)']])


# Extract Movie Names and Durations
Movie_Name_duration = top_10_duration['Movie_Name']
Duration = top_10_duration['Duration(Minutes)']
Rank_duration = top_10_duration['Rank']

y_pos = np.arange(len(Movie_Name_duration))

# Create the horizontal bar chart
plt.barh(y_pos, Duration, align='center')
plt.yticks(y_pos, [f'{r}. {m}' for r, m in zip(Rank_duration, Movie_Name_duration)])
plt.xlabel('Duration (Minutes)')
plt.title('Top 10 Movies by Duration')

plt.show()
```

	Rank	Movie_Name	Duration(Minutes)
159	1	Gone with the Wind	238
82	2	Once Upon a Time in America	229
97	3	Lawrence of Arabia	218
184	4	Ben-Hur	212
21	5	Shichinin No Samurai	207
3	6	The Godfather: Part II	202
6	7	The Lord of the Rings: The Return of the King	201
5	8	Schindler's List	195
27	9	The Green Mile	189
186	10	Barry Lyndon	185



```
In [66]: min_yr = df['Released_Year'].min()
print(min_yr)

1921
```

```
In [67]: max_yr = df['Released_Year'].max()
print(max_yr)

2023
```


```
In [68]: df['Decade'] = (df['Released_Year'] // 10) * 10
```

```
# Calculate the average ratings for each decade
average_ratings_by_decade = df.groupby('Decade')['Ratings'].mean()

# Plot the average ratings in a line chart
plt.plot(average_ratings_by_decade.index, average_ratings_by_decade.values, marker='o', linestyle='-')


# Set labels and title
plt.xlabel('Decade')
plt.ylabel('Average Ratings')
plt.title('Average Ratings by Decade')

# Show the plot
plt.show()
```



```
In [69]: import seaborn as sns
sns.histplot(x=df['Duration(Minutes)'])
plt.title('Movie Duration Distribution')
plt.ylabel('Movies Count')
```

```
Out[69]: Text(0, 0.5, 'Movies Count')
```




```
In [73]: from wordcloud import WordCloud
high_rated_movies = df[df['Ratings'] > 8.0]

# Create a string containing movie names and their ratings
movie_ratings_text = ' '.join(f"{name} {rating}" for name, rating in zip(high_rated_movies['Movie_Name'], high_rated_movies['Ratings']))


# Generate the word cloud
wordcloud = WordCloud(width=800, height=400, background_color='white').generate(movie_ratings_text)

# Display the word cloud
plt.figure(figsize=(10, 5))
plt.imshow(wordcloud, interpolation='bilinear')
plt.axis('off')
plt.show()
```



```
In [70]: sns.scatterplot(x=df['Votes'],y=df['Ratings'])
plt.title('Correlation between Ratings and Votes')
```

```
Out[70]: Text(0.5, 1.0, 'Correlation between Ratings and Votes')
```



```
In [75]: Q1 = df['Votes'].quantile(0.25)
Q3 = df['Votes'].quantile(0.75)
IQR = Q3 - Q1

# Define the lower and upper bounds for outliers
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR

# Find outliers
outliers = df[(df['Votes'] < lower_bound) | (df['Votes'] > upper_bound)]

print("Outliers:")
print(outliers)
```

```
Outliers:
   Rank  Movie_Name  Duration(Minutes)  Released_Year \
0      1  The Shawshank Redemption          142         1994
2      3      The Dark Knight          152         2008
7      8      Pulp Fiction          154         1994
10     11  Forrest Gump          142         1994
11     12  Fight Club          139         1999
13     14  Inception          148         2010
```

```

   Certification  Ratings  Votes  Decade
0              A      9.3  2800000   1990
2              UA      9.0  2800000   2000
7              A      8.9  2200000   1990
10             UA      8.8  2200000   1990
11              A      8.8  2200000   1990
13             UA      8.8  2500000   2010
```

```
In [76]: from sklearn.linear_model import LinearRegression
x = df[['Votes']]

# Fit a linear regression model
model = LinearRegression()
model.fit(X, df['Votes'])

# Predict 'Votes' using the linear regression model
predicted_votes = model.predict(X)

# Calculate the residuals
residuals = df['Votes'] - predicted_votes
```

```
# Identify potential outliers (e.g., points with residuals > 2 standard deviations)
outliers = df[np.abs(residuals) > 2 * np.std(residuals)]
```

```
# Print the potential outliers
print("Potential Outliers:")
print(outliers)
```

Potential Outliers:

Empty DataFrame

Columns: [Rank, Movie_Name, Duration(Minutes), Released_Year, Certification, Ratings, Votes, Decade]

Index: []

```
In [77]: highest_rating_movie = df[df['Ratings'] == df['Ratings'].max()]

# Find the movie with the maximum votes
max_votes_movie = df[df['Votes'] == df['Votes'].max()]

# Find the movie with the maximum duration
max_duration_movie = df[df['Duration(Minutes)'] == df['Duration(Minutes)'].max()]

# Combine all the relevant information
combined_info = pd.concat([highest_rating_movie, max_votes_movie, max_duration_movie], ignore_index=True)

# Create a figure to display the information
fig, ax = plt.subplots(figsize=(10, 5))

# Plot the information
for index, row in combined_info.iterrows():
    ax.text(0.1, 0.9 - index*0.2, f"Movie Name: {row['Movie_Name']}\nRatings: {row['Ratings']}\nVotes: {row['Votes']}\nDuration: {row['Duration(Minutes)']}",
           fontsize=12, ha='left', va='center', transform=ax.transAxes)

# Remove the axes for a cleaner look
ax.axis('off')

plt.show()
```

Movie Name: The Shawshank Redemption

Ratings: 9.3

Votes: 2800000

Duration: 142 minutes

Movie Name: The Shawshank Redemption

Ratings: 9.3

Votes: 2800000

Duration: 142 minutes

Movie Name: The Dark Knight

Ratings: 9.0

Votes: 2800000

Duration: 152 minutes

Movie Name: Gone with the Wind

Ratings: 8.2

Votes: 327000

Duration: 238 minutes

```
In [79]: highest_rating_movie = df[df['Ratings'] == df['Ratings'].max()]

# Find the movie with the maximum votes
max_votes_movie = df[df['Votes'] == df['Votes'].max()]

# Find the movie with the maximum duration
max_duration_movie = df[df['Duration(Minutes)'] == df['Duration(Minutes)'].max()]

# Combine all the relevant information
combined_info = pd.concat([highest_rating_movie, max_votes_movie, max_duration_movie], ignore_index=True)

# Create a bar chart to display the information
fig, ax = plt.subplots(figsize=(10, 5))

# Plot the information
ax.barh(combined_info['Movie_Name'], combined_info['Ratings'], color='blue', label='Ratings')
ax.barh(combined_info['Movie_Name'], combined_info['Votes'], color='green', label='Votes')
ax.barh(combined_info['Movie_Name'], combined_info['Duration(Minutes)'], color='red', label='Duration')

# Add labels and legend
ax.set_xlabel('Values')
ax.set_title('Highest votes, rating, and duration movie')
ax.legend()

plt.show()
```

Highest votes, rating, and duration movie

